

Introduction to Computer Graphics

2016 Spring

National Cheng Kung University

Instructors: Min-Chun Hu 胡敏君

Shih-Chin Weng 翁士欽 (西基電腦動畫)




Buffers and Mapping Techniques



Buffer

- Define a buffer as a block of memory with $n \times m$ elements in spatial resolution and k elements in depth/precision.
 - k : the number of bits used to represent each pixel



OpenGL Frame Buffers

- Each frame buffer can have different depth (k) and can be represented in different data type


- 64 bits RGBA Color Buffers

- 32 bits Front Buffer (byte)

- 32 bits Back Buffer (byte)


- 32 bits Depth Buffer (integer or floating point)

- ...



Writing in Buffers


- Conceptually, we can consider all of memory as a large two-dimensional array of pixels
- We only occasionally read and write a rectangular block of pixels (i.e. frame buffer) inside the memory
 - Bit block transfer (bitblt) operations or raster operations (raster-ops)



`write_block(source, n, m, x, y, destination, u, v)`

Writing Model


- Read destination pixel before writing source



Bit Writing Modes

- Source and destination bits are combined bitwise
- 16 possible functions/writing modes (one per column in table)

		<div>clear replace XOR OR clear</div>															
s	d	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1
0	1	0	0	0	0	1	1	1	1	0	0	0	0	1	1	1	1
1	0	0	0	1	1	0	0	1	1	0	0	1	1	0	0	1	1
1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1



```
glLogicOp(mode);
glEnable(GL_COLOR_LOGIC_OP);
```

The Limits of Geometric Modeling


- Although graphics cards can render over 10 million polygons per second, the number is insufficient for many phenomena
 - Clouds
 - Grass
 - Terrain
 - Skin



Image from “Final Fantasy” movie

Modeling an Orange

- Consider the problem of modeling an orange
- Start with an orange-colored sphere
 - Too simple
- Replace sphere with a more complex shape
 - Does not capture surface characteristics (small dimples)
 - Takes too many polygons to model all the dimples




Mapping


- Instead, we use mapping method that build a simple model and add details as part of the rendering process.
- Three major mapping methods:
 - Texture Mapping
 - Environment (Reflection) Mapping
 - Bump Mapping

Texture Mapping

- Use an image (or texture) to fill inside of polygons
 - Take a picture of a real orange, scan it, and “paste” onto simple geometric model
- Still might not be sufficient because resulting surface will be smooth
 - Need to change local shape
 - Bump mapping



Texture Mapping



geometric model



texture mapped


Environment (Reflection) Mapping

- Allows simulation of highly specular surfaces
- An image of the environment is painted onto the surface as that surface is being rendered
 - Create images that have the appearance of reflected materials without our having to trace reflected rays




Bump Mapping

- Distort the normal vectors during the rendering process to make the surface appear to have small variations in shapes




The Graphics Pipeline



- The graphics pipeline or rendering pipeline refers to the sequence of steps used to create a 2D raster representation of a 3D scene/model.
- Vertex processing
 - Each vertex is processed independently.
 - To carry out coordinate transformations.
 - Each change of the camera coordinate can be represented by a matrix.
 - To compute a color for each vertex.
- Clipper and Primitive Assembly
 - Efficient clipping must be done on a primitive-by-primitive basis rather than on a vertex-by-vertex basis.

The Graphics Pipeline (Cont.)



■ Rasterization (Scan conversion)

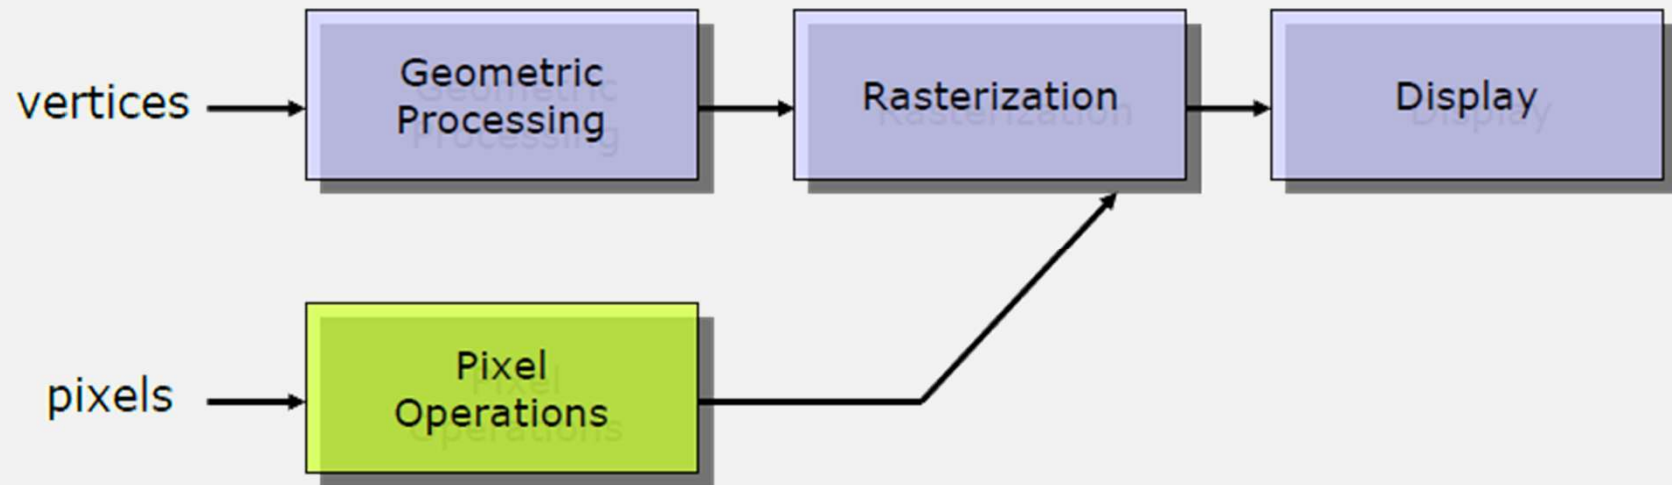
- Primitives emerging from the clipper are still represented in terms of their vertices and must be converted to pixels in the frame buffer.
- Determine which pixels in the frame buffer are inside the polygon.
- Output of rasterization is a set of **fragments** (potential pixels with color, location, and depth information) for each primitive.

■ Fragment Processing

- Update the pixels in the frame buffer according to the processed fragments. (Some surfaces may not be visible because of occlusion)
- The color of pixels in each fragment can be altered by **texture mapping** or **bump mapping**.

Where Does Mapping Take Place?

- Mapping techniques are implemented at the end of the rendering pipeline



Two-Dimensional Texture Mapping in OpenGL

■ Three basic steps:

■ 1st step: form a texture image and place it in texture memory on the GPU

▣ `glTexImage2D(Glenum target, GLint level, GLint iformat, GLsizei width, GLsizei height, GLint border, GLenum format, GLenum type, GLvoid *tarray);` // specify a two-dimensional texture

- target : choose a single image, set up a cube map, or test if there is sufficient texture memory
- level: used for mipmapping, where 0 denotes the highest level or we are not using mimapping
- iformat: specifies how we would like the texture stored in texture memory
- width/height: specify the size of the image in the memory
- format/type: describe how pixels in the image in processor memory are stored, so that OpenGL can read those pixels and store them in texture memory

■ 2nd step: assign texture coordinates to each fragment


■ 3rd step: apply the texture to each fragment

■ Multiple ways to accomplish each step

■ Controlled by many parameters

Two-Dimensional Texture Mapping in OpenGL (Cont.)

- Use two floating-point texture coordinates, s and t , rather than using integer texel locations that depend on the dimension of texture image



Texture Mapping Example


- The texture (below) is a 256 x 256 image, mapped to a rectangular polygon which is viewed in perspective.



Is it Simple?

- Although the idea is simple

- there are 3 or 4 coordinate systems involved in mapping an image to a surface



Coordinate Systems

- Parametric coordinates

- Used to model curves and surfaces

- Object or World Coordinates

- Conceptually, where the mapping takes place

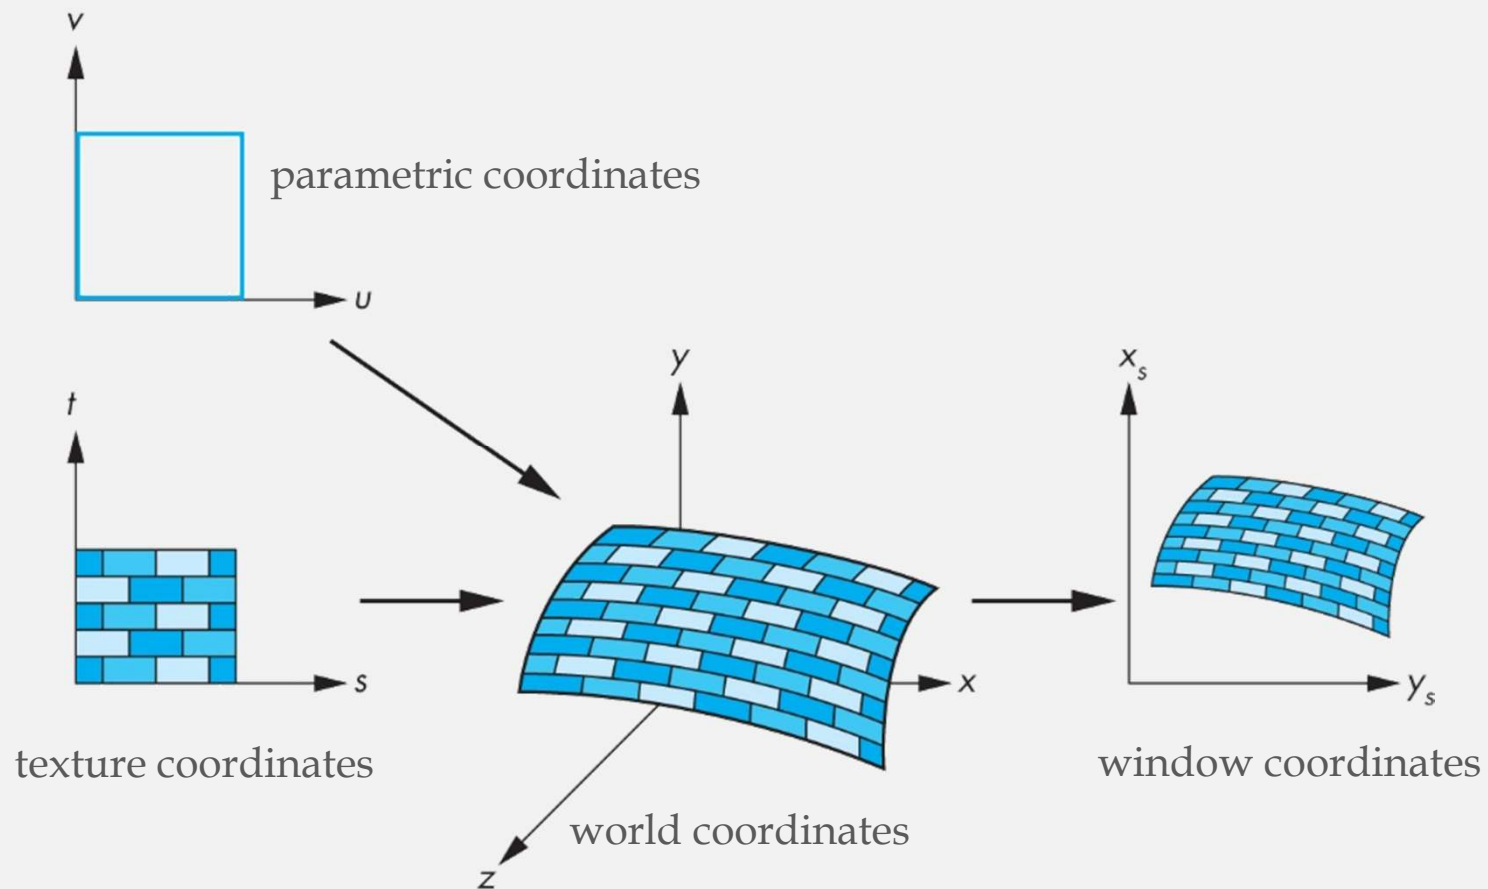
- Texture coordinates

- Used to identify points in the image to be mapped

- Window Coordinates

- Where the final image is really produced

Texture Mapping Involving 3 Mappings



Mapping Functions


- Basic problem is how to find the maps
- Consider mapping from texture coordinates to a point a surface
- Appear to need three functions

■ $x = x(s,t)$

■ $y = y(s,t)$

■ $z = z(s,t)$

■ $w = w(s,t)$




Backward Mapping

- Given a point (x, y, z) or (x, y, z, w) on an object, find the corresponding texture coordinates, i.e. the texel $T(s,t)$
 - $s = s(x, y, z, w)$
 - $t = t(x, y, z, w)$
- Such functions are difficult to find in general

Linear Map

- Map a point in the texture map $T(s,t)$ to a point on the surface $p(u,v)$ by a linear map



$$u = as + bt + c$$

$$v = ds + et + f$$

$$u = u_{min} + \frac{s - s_{min}}{s_{max} - s_{min}} (u_{max} - u_{min})$$


$$v = v_{min} + \frac{s - s_{min}}{s_{max} - s_{min}} (v_{max} - v_{min})$$

- Does not take into account the curvature of the surface

Two-part Mapping

■ 1st step: Map the texture to a simple intermediate surface

■ Example: map to cylinder, sphere, or cube



parametric cylinder:

$$x = r \cos 2\pi u$$

$$y = r \sin 2\pi u$$

$$z = v/h$$


■ 2nd step: Map the intermediate surface to the surface being rendered

Spherical Map

- We can use a parametric sphere
 - $x = r \cos 2\pi u$
 - $y = r \sin 2\pi u \cos 2\pi v$
 - $z = r \sin 2\pi u \sin 2\pi v$
- in a similar manner to the cylinder but have to decide where to put the distortion
- Spheres are used in environmental maps


Box Mapping

- Easy to use with simple orthographic projection
- Also used in environment maps (Cube mapping)




Second Mapping

- Map from an intermediate object to an actual object
 - Normals from intermediate to actual
 - Normals from actual to intermediate
 - Vectors from center of intermediate




Two-part Mapping



Texture Sampling

■ How to assign a texture value to a pixel ?

- Point Sampling: use the value of the texel that is closest to the texture coordinate output by the rasterizer
- Linear Filtering: weighted averaging the neighborhood texels of the texel determined by point sampling
 - A better strategy
 - More difficult to implement (how to deal with the **boundary of the texel array**)
 - Still imperfect due to the limited **resolution of both the frame buffer and the texture map**



Magnification and Minification



- The size of the pixel on the screen may be smaller or larger than one texel

- Magnification

- The texel covers multiple pixels


- Minification

- The pixel covers multiple texels




Magnification

- The alignment is probably not exact.




Nearest Texel

- Find the nearest texel.



Nearest Texel

- Find the nearest texel.




Linear Interpolation

- OpenGL may also interpolate the colors of the nearest four texels.


Linear Interpolation

- Find the nearest four texels.




Linear Interpolation

- Find the nearest four texels.



Example: Interpolation

- Using the nearest texel, color the pixels.



Example: Interpolation

- Compute the color of the pixel (2, 4).
- Assume the texture is 2×2 .
- The center of the pixel is
 - 25% of the way across the group of texels.
 - Therefore, $s = 0.25$.
 - 50% of the way up the group of texels.
 - Therefore, $t = 0.50$.

Example: Interpolation

■ Interpolate horizontally:

■ Top edge:

$$0.75(1, 0, 0) + 0.25(0, 1, 0) = (0.75, 0.25, 0).$$

■ Bottom edge:

$$0.75(0, 0, 1) + 0.25(1, 1, 0) = (0.25, 0.25, 0.75).$$


■ Now interpolate those values vertically:

$$\begin{aligned} & \text{■ } 0.5(0.75, 0.25, 0) + 0.5(0.25, 0.25, 0.75) \\ & = (0.5, 0.25, 0.375). \end{aligned}$$




Minification

- Again, the alignment is not exact.



Minification

- If 64 texels all map to the single pixel, what color should the pixel be?




Minification

- Again, we may choose between the nearest texel and interpolating among the nearest four texels.

Minification

- Choose the nearest texel.




Minification


- If we choose to interpolate, then OpenGL will compute the average of the four texels that whose centers are *nearest* to the center of the pixel.
- This will reduce, but not eliminate, the aliasing or effect.

Minification

- Choose the four nearest texels.




Minification



Aliasing


- Point sampling of the texture can lead to aliasing errors



Aliasing Example



Original image




Sample one for each 5x5 pixels

Ref: www.relisoft.com/Science/Graphics/alias.html

Area Averaging

- A better but slower option is to use area averaging




Area Averaging



Original image



Applying a 5x5 box filter



Sampling every 5x5 pixels




Sampling every 5x5 pixels

[www.relisoft.com/
Science/Graphics/
alias.html](http://www.relisoft.com/Science/Graphics/alias.html)

Mipmapped Textures



- OpenGL has another way to deal with the minification problem:
 - Mipmapping: create a series of texture arrays at reduced sizes
 - `glGenerateMipmap(GL_TEXTURE_2D);`



storage


Mipmapping

- 1/3 overhead of maintaining the MIP map.




Example

point sampling




linear filtering





mipmapped
point sampling



mipmapped
linear filtering




Mesh Size Problem



Texture Mapping for Polygons in OpenGL

- Based on parametric texture coordinates

- `glTexCoord*()` specified at each vertex



Interpolation

- OpenGL uses interpolation to find proper texels from specified texture coordinates

- Can be distortions


good selection
of tex coordinates



poor selection
of tex coordinates





texture stretched
over trapezoid
showing effects of
bilinear interpolation



Interpolation (Cont.)



- Can we just use Linear interpolation in screen space?



Pictures from lecture notes of “Computer Graphics”, UNC

Reduction of the flaws

- Subdivide the texture-mapped triangles into smaller triangles.



- Is it correct?

Mapping from Screen Space to 3D Space


■ Interpolation in screen space

$$\blacksquare P(m) = P_1 + m(P_2 - P_1)$$

■ Interpolation in 3D space

$$\blacksquare V(n) = V_1 + n(V_2 - V_1)$$

$$\blacksquare P_y(n) = V_y(n) / V_z(n)$$



$$P_y(m) = \frac{y_1}{z_1} + m \left(\frac{y_2}{z_2} - \frac{y_1}{z_1} \right) = \frac{y_1 + n(y_2 - y_1)}{z_1 + n(z_2 - z_1)} \Rightarrow n = \frac{mz_1}{z_2 + m(z_1 - z_2)}$$

Perspective Correct Interpolation

- $T(n) = T_1 + n(T_2 - T_1)$

- Assume $w_1 = \frac{1}{z_1}$, $w_2 = \frac{1}{z_2}$ (for the graphics pipeline)

$$\begin{aligned} I &= I_1 + \frac{mz_1}{z_2 + m(z_1 - z_2)} (I_2 - I_1) \\ &= I_1 + \frac{mw_2}{w_1 + m(w_2 - w_1)} (I_2 - I_1) \\ &= \frac{I_1w_1 + m(I_2w_2 - I_1w_1)}{w_1 + m(w_2 - w_1)} \end{aligned}$$

How to Handle Highly Specular Surfaces?

- How to render a flat mirror?
- How to render a mirror-like object in a virtual scene?




Environment Mapping

- Also known as reflection mapping
- First proposed by Blinn and Newell.
- An efficient way to create reflections on curved surfaces
 - can be implemented using texture mapping supported by graphics hardware


Environment Mapping (Cont.)

- Assume the environment is far away and there's no self-reflection
- The reflection at a point can be solely decided by the reflection vector.




Environment Mapping (Cont.)

- We can obtain an approximately correct value of the shade as part of a two-step rendering pass
 - 1st step: render the scene without mirror polygon, with the camera placed at the center of the mirror pointed in the direction of the normal of the mirror
 - Thus, we obtained an image of objects seen by the mirror
 - 2nd step: render the scene with the mirror polygon, and use the obtained image as the shade (texture) to place on the mirror polygon




Environment Mapping (Cont.)

- Project the environment onto a sphere centered at the center of projection




Environment Mapping Pictures



Pictures from lecture notes of Computer Graphics course, UNC


Sphere Mapping

- The image texture is taken from a perfectly reflective sphere.
- Assume the size of the sphere $\rightarrow 0$.
 - Map the rays to the environment




Pictures from OpenGL tutorial. <http://www.opengl.org>

Sphere Mapping (Cont.)



Sphere Mapping (Cont.)



■ To access the sphere map texture

■ Compute the reflection vector on the object surface as usual

$$r = (r_x, r_y, r_z) = e' - 2(n' \cdot e')n'$$

■ Now, compute the sphere normal in the local space


$$n'' = (r_x, r_y, r_z) + (0, 0, -1)$$

$$n''' = \left(\frac{r_x}{m}, \frac{r_y}{m}, \frac{r_z - 1}{m} \right) \quad m = \sqrt{r_x^2 + r_y^2 + (r_z - 1)^2}$$



■ Normalized the screen space from [-1,1] to [0,1]

$$s = \frac{r_x}{2m} + \frac{1}{2}$$

$$t = \frac{r_y}{2m} + \frac{1}{2}$$




Reflective Cube Map



Bump and Normal Mapping

- Represent surface details and avoid heavy geometric computation.



http://www.ozone3d.net/tutorials/bump_mapping.php

Bump and Normal Mapping (Cont.)

- Calculate reflection (Phong Shading) with a normal map.
- Or with a height map.

Smooth surface



Bumpy surface



Bump-mapped
surface



Bump Mapping

- Let $\mathbf{p} = \mathbf{p}(u, v)$ be a smooth parametric surface, with normals $\mathbf{n} = \mathbf{n}(u, v)$.

$$\mathbf{p}_u = \begin{bmatrix} \frac{\partial x}{\partial u} \\ \frac{\partial y}{\partial u} \\ \frac{\partial z}{\partial u} \end{bmatrix} \quad \mathbf{p}_v = \begin{bmatrix} \frac{\partial x}{\partial v} \\ \frac{\partial y}{\partial v} \\ \frac{\partial z}{\partial v} \end{bmatrix} \quad \mathbf{n} = \mathbf{p}_u \times \mathbf{p}_v$$

- Apply a bump map $b = b(u, v)$:

$$\mathbf{p}' = \mathbf{p} + d(u, v)\mathbf{n} \quad \text{Displaced surface}$$

$$\mathbf{n}' = \mathbf{p}_u' \times \mathbf{p}_v' \quad \text{Normal at } \mathbf{p}'$$


$$\mathbf{p}_u' = \frac{\partial}{\partial u} (\mathbf{p} + d(u, v)\mathbf{n}) = \mathbf{p}_u + d_u\mathbf{n} + d(u, v)\mathbf{n}_u$$

$$\mathbf{p}_v' = \frac{\partial}{\partial v} (\mathbf{p} + d(u, v)\mathbf{n}) = \mathbf{p}_v + d_v\mathbf{n} + d(u, v)\mathbf{n}_v$$

$$\mathbf{n}' \approx \mathbf{n} + d_u(\mathbf{n} \times \mathbf{p}_v) + d_v(\mathbf{n} \times \mathbf{p}_u)$$


Opacity and Transparency

- Opaque surfaces permit no light to pass through
- Transparent surfaces permit all light to pass
- Translucent surfaces pass some light
 - translucency = $1 - \text{opacity } (\alpha)$



Physical Models

- Dealing with translucency in a physically correct manner is difficult due to
 - the complexity of the internal interactions of light and matter
 - Using a pipeline renderer



Blending Equation

- We can define source and destination blending factors for each RGBA component

- $s = [s_r, s_g, s_b, s_\alpha]$

- $d = [d_r, d_g, d_b, d_\alpha]$

- Suppose that the source and destination colors are

- $b = [b_r, b_g, b_b, b_\alpha]$

- $c = [c_r, c_g, c_b, c_\alpha]$

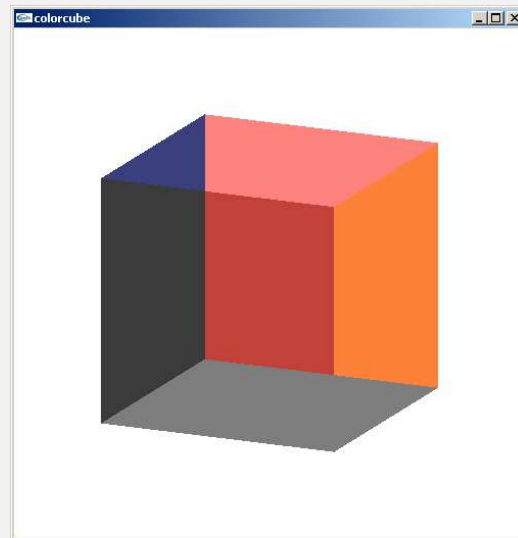
- Blend as

- $c' = [b_r s_r + c_r d_r, b_g s_g + c_g d_g, b_b s_b + c_b d_b, b_\alpha s_\alpha + c_\alpha d_\alpha]$

Order Dependency

■ Is this image correct?


- Probably not
- Polygons are rendered in the order they pass down the pipeline
- Blending functions are order dependent



Fog

- We can composite with a fixed color and have the blending factors depend on depth
 - Simulates a fog effect
 - Blend source color C_s and fog color C_f by
 - $C_s' = f C_s + (1 - f)C_f$
- f is the fog factor
 - Exponential
 - Gaussian
 - Linear

Fog Functions




Accumulation Buffer

- Compositing and blending are limited by resolution of the frame buffer
 - Typically 8 bits per color component
- The accumulation buffer is a high resolution buffer
 - 16 or more bits per component
 - Write into it or read from it with a scale factor
- Slower than direct compositing into the frame buffer


Applications of Composition Techniques

- Compositing
- Image Filtering
- Whole scene antialiasing
- Motion effects
-

Depth of Focus



Motion blur



http://www.eml.hiroshima-u.ac.jp/gallery/ComputerGraphics/motion_blur/